

**Root-Cause Analysis through Reinforcement  
Learning  
To  
AFLRL (AMERICAN FUZZY LOP FOR  
REINFORCEMENT LEARNING)**

*By KP*

# Overview


- RCA: For finding security vulnerabilities through fuzzing.
- RCA is slow, I mean very slow.
- Proper mutation for sample needed.
- Reinforcement Learning Implementation.
- Counterexamples are the key.
- Approach of using RACING increasing scalability and better than today's RCA.

# Introduction

- Need for automation for faster devops cycle.
- CVE-2017-5380, CVE-2018-4145 and CVE-2022-36320.
- Try to find the crash is the key
- How to achieve that? Through mutation of test cases.
- Root cause analysis which has been done manually.

SLow SLow Slowwwwww.....

# RCA



# Example

## An example: CVE-2019-9077

```
16182 eopt = get_data (NULL, filedata, options_offset, 1,
16183                sect->sh_size, _("options"));
16184 if (eopt)
16185 {
16186     iopt = (Elf_Internal_Options *)
16187         calloc ((sect->sh_size / sizeof (eopt)), sizeof (* iopt));
16188     ...
16194     offset = 0;
16195     option = iopt;
16196     while (offset <= sect->sh_size - sizeof (* eopt))
16197     {
16199         Elf_External_Options * option;
16200
16201         eoption = (Elf_External_Options *)((char *) eopt + offset);
16202
16203         option->kind = BYTE_GET (eoption->kind);
16204         option->size = BYTE_GET (eoption->size);
16205
16206         offset += option->size;
16207         ++option;
16208     }
16209     ...

```

sect->sh\_size=1 < 8

sect->sh\_size=1 < 8

Root Cause

Program crashes


# Automated RCA Challenges

- Having a pipeline with bucket of analysis tools.
- They rely on rules and types for finding bug.
- Lead to type-confusion flaws.


# Spectrum Based Fault Localization (SFL)

- Finding bug without explicit data dependency.
- They outperform traditional automated RCA.
- They try to correlate crash and root cause through statistics.
- Ranking highly correlated test case and find the crash.
- Takes lot of time 12 hours or even week.

# Statistical RCA



```
16197 while (offset <= sect->sh_size - sizeof (* eopt))
16198 {
16199     Elf_External_Options * option;
16200     eoption = (Elf_External_Options *)((char *) eopt + offset)
16201
16202     option->kind = BYTE_GET (eoption->kind);
16203     option->size = BYTE_GET (eoption->size);
16204
16205     offset += option->size;
16216     ++option;
16217 }
16218 }
```






# General Workflow of Statistical RCA



**Test Case  
Collection**



**Predicate Construction and  
Correlation Analysis**






**Ranking**

# Racing

What is Racing? Root-cause-analysis on Counter-examples based reinforcement-learning.....

# How to mutate test case?




AFL Fuzzer

Seed  
Selection

Mutation  
Selection

action

reward



Reinforcement  
learning

# All about Correlation

predicate	correlation
p <sub>1</sub> : x>10	1.00
p <sub>2</sub> : y<8	0.95 ↓
p <sub>3</sub> : size<32	0.93
p <sub>4</sub> : i>0	0.90

CoP<sub>1</sub>  
←


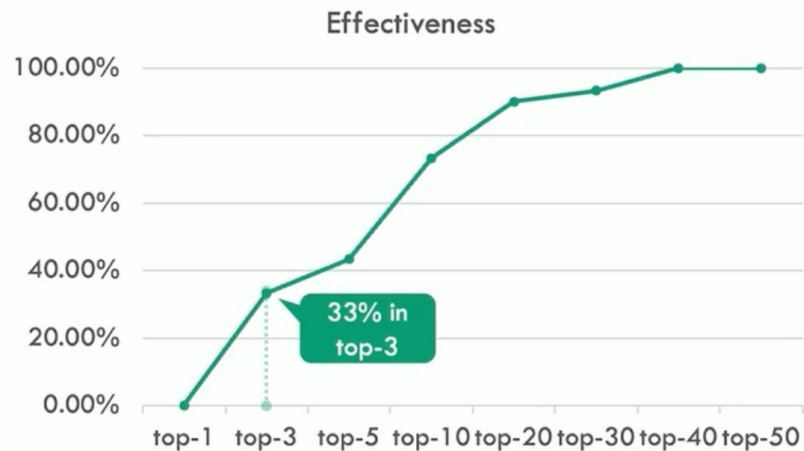
predicate	correlation
p <sub>1</sub> : x>10	1.00
p <sub>2</sub> : y<8	0.97
p <sub>3</sub> : size<32	0.93
p <sub>4</sub> : i>0	0.90

CoP<sub>2</sub>  
→

predicate	correlation
p <sub>1</sub> : x>10	1.00
p <sub>3</sub> : size<32	0.93
p <sub>2</sub> : y<8	0.91 ↓
p <sub>4</sub> : i>0	0.90

$$g_t = g_t^{\text{count}} + g_t^{\text{order}}$$

# Existing Results



# Algorithm code

```
1 # RACING Algorithm Template
2
3 # Initialize the algorithm state
4 S_t = [] # Initial state
5 gamma_t = 0.5 # Exploration-Exploitation tradeoff parameter
6 t = 1 # Iteration counter
7
8 # Function placeholders (to be implemented based on your needs)
9 def select_action(S_prev):
10     """Select an action based on the previous state."""
11     pass
12
13 def sample_action(d_t, gamma_t):
14     """Sample an action based on the decision and gamma parameter."""
15     pass
16
17 def generate_input(action):
18     """Generate a test input based on the selected action."""
19     pass
20
21 def run_program(prog, inputs):
22     """Run the program with the given inputs and return execution results."""
23     pass
24
25 def rank_results(results):
26     """Rank the predicates based on the execution results."""
27     pass
28
29 def compute_reward(results, S_prev):
30     """Compute the reward for the generated inputs and state."""
31     pass
32
33 def update_state(S_prev, rankings, reward, results):
34     """Update the state based on new information."""
35     pass
```

```
def update_gamma(rewards, decisions, actions):
    """Update the gamma parameter based on rewards and actions."""
    pass

def has_converged(state):
    """Check if the algorithm has converged."""
    pass

# Main loop
while True:
    d_t = select_action(S_t) # Select an action
    a_t = sample_action(d_t, gamma_t) # Sample an action
    inputs = generate_input(a_t) # Generate inputs
    results = run_program("Prog", inputs) # Run the program
    rankings = rank_results(results) # Rank the results
    reward = compute_reward(results, S_t) # Compute the reward
    S_t = update_state(S_t, rankings, reward, results) # Update state
    gamma_t = update_gamma(reward, d_t, a_t) # Update gamma parameter

    if has_converged(S_t): # Check convergence
        break

# Final step: Select top-50 predicates
top_50_predicates = sorted(rankings, key=lambda x: x[1], reverse=True)[:50]
r_star = [predicate for predicate, score in top_50_predicates]

# Output results
print("Top-50 Predicates:", top_50_predicates)
print("True Order (r*):", r_star)
```

# Optimization in a code

optimized\_racing.py > ...

```
1 import random
2 import numpy as np
3 from collections import defaultdict
4
5 # Initialize state and parameters
6 S_t = [] # Initial state
7 gamma_t = 0.5 # Exploration-Exploitation tradeoff parameter
8 t = 1 # Iteration counter
9 MAX_ITERATIONS = 100 # Maximum iterations to avoid infinite loops
10 CONVERGENCE_THRESHOLD = 0.01 # Threshold for convergence
11 top_k = 50 # Number of top predicates to select
12
13 # Helper functions for optimization
14 def select_action(S_prev):
15     """
16     Select an action based on the previous state.
17     Here, a simple weighted random selection is implemented. Replace with a more sophisticated
18     """
19     if not S_prev:
20         return random.choice(["Action1", "Action2", "Action3"])
21     # Weighted selection based on state information
22     return random.choices(["Action1", "Action2", "Action3"], weights=[0.3, 0.5, 0.2])[0]
23
24 def sample_action(d_t, gamma_t):
25     """
26     Sample an action by balancing exploration and exploitation.
27     """
28     if random.random() < gamma_t:
29         return f"Explore_{d_t}" # Exploration
30     return f"Exploit_{d_t}" # Exploitation
```

```
def generate_input(action):
    """
    Generate a new test input based on the action.
    """
    # Simple mutation logic. Replace with advanced input generation if necessary.
    return f"Input_based_on_{action}"

def run_program(prog, inputs):
    """
    Run the program with the given inputs.
    Mock execution results here for demonstration purposes.
    """
    results = [{"predicate": f"Predicate_{i}", "value": random.random()} for i in range(10)]
    return results

def rank_results(results):
    """
    Rank predicates based on their impact (value).
    """
    return sorted(results, key=lambda x: x["value"], reverse=True)

def compute_reward(results, S_prev):
    """
    Compute a reward for the current iteration based on improvements.
    Reward is based on discovering new predicates or improving ranks.
    """
    new_coverage = sum([result["value"] for result in results])
    old_coverage = sum([state["value"] for state in S_prev]) if S_prev else 0
    return new_coverage - old_coverage
```

# Continuation

```
def update_state(S_prev, rankings, reward, results):  
    """  
    Update the state based on new rankings and results.  
    """  
    updated_state = S_prev + results  
    # Keep only unique predicates  
    unique_state = {result["predicate"]: result for result in updated_state}  
    return list(unique_state.values())
```

```
def update_gamma(rewards, decisions, actions):  
    """  
    Update the gamma parameter dynamically based on the rewards.  
    """  
    if rewards > 0:  
        return min(1.0, gamma_t + 0.05) # Favor exploitation  
    return max(0.1, gamma_t - 0.05)      # Favor exploration
```

```
def has_converged(state, prev_state):  
    """  
    Check if the state has converged.  
    """  
    if not prev_state:  
        return False  
    diff = sum(abs(curr["value"] - prev["value"]) for curr, prev in zip(state, prev_state))  
    return diff < CONVERGENCE_THRESHOLD
```

```
# Main optimization loop  
prev_state = []  
while t <= MAX_ITERATIONS:  
    d_t = select_action(S_t)           # Select an action  
    a_t = sample_action(d_t, gamma_t)  # Sample an action  
    inputs = generate_input(a_t)       # Generate inputs  
    results = run_program("Prog", inputs) # Run the program  
    rankings = rank_results(results)    # Rank predicates  
    reward = compute_reward(results, S_t) # Compute the reward  
    prev_state = S_t  
    S_t = update_state(S_t, rankings, reward, results) # Update state  
    gamma_t = update_gamma(reward, d_t, a_t)           # Update gamma parameter  
  
    if has_converged(S_t, prev_state):                 # Check convergence  
        print(f"Converged after {t} iterations.")  
        break  
    t += 1
```

```
# Final step: Select top-k predicates  
top_predicates = rank_results(S_t)[:top_k]  
print(f"Top-{top_k} Predicates: {top_predicates}")
```



# Unexpected Turn

- The original code base from authors were not producing proper results.

```
clang++ llvm-config --cxxflags -Wl,-znodelete -fno-rtti -fpic -O3 -funroll-loops -Wall -D_FORTIFY_
riadic-macros -shared full_trace.cpp afl-llvm-pass.so.cc -o ../afl-llvm-pass.so `llvm-config --ldfl
full_trace.cpp:28:10: fatal error: 'llvm/IR/CallSite.h' file not found
 28 | #include "llvm/IR/CallSite.h"
    |
1 error generated.
afl-llvm-pass.so.cc:64:10: fatal error: 'llvm/Transforms/IPO/PassManagerBuilder.h' file not found
 64 | #include "llvm/Transforms/IPO/PassManagerBuilder.h"
    |
1 error generated.
./02_PocExecutionInspector.sh: line 9: cd: /Racing-eval/scripts/: No such file or directory
python3: can't open file '/home/ubuntu/Racing-code/examples/1-readelf-heap-buffer-overflow/tracing.py': [Errno 2] No such file or dire
python3: can't open file '/home/ubuntu/Racing-code/examples/1-readelf-heap-buffer-overflow/binutils-2.32/binutils/readelf': [Errno 2] No such file or dire
configure: error: in `/home/ubuntu/Racing-code/examples/1-readelf-heap-buffer-overflow/binutils-2.32':
configure: error: C compiler cannot create executables
See `config.log' for more details
make: *** No targets specified and no makefile found. Stop.
mv: cannot stat '/home/ubuntu/Racing-code/examples/1-readelf-heap-buffer-overflow/binutils-2.32/binutils/readelf': No such file or direc
rm -rf texinfo/doc/Makefile texinfo/po/POTFILES
rmdir texinfo/doc texinfo/info texinfo/intl texinfo/lib 2>/dev/null
make: [Makefile:2069: local-distclean] Error 1 (ignored)
rmdir texinfo/makeinfo texinfo/po texinfo/util 2>/dev/null
make: [Makefile:2070: local-distclean] Error 1 (ignored)
rmdir fastjar gcc gnattools gotools libcc1 libiberty 2>/dev/null
make: [Makefile:2071: local-distclean] Error 1 (ignored)
rmdir texinfo zlib 2>/dev/null
make: [Makefile:2072: local-distclean] Error 1 (ignored)
find . -name config.cache -exec rm -f {} \; \; 2>/dev/null
make: [Makefile:2073: local-distclean] Error 1 (ignored)
```

- Because of these errors and file issues, I changed the direction of my project.

# Building AFLRL

New Focus:

- **"AFLRL"**: Training reinforcement learning models using AFL (American Fuzzy Lop).
- Build datasets for RL model training.
- Enhance AFL's capabilities by adding vulnerable code paths to generate robust training data.

# Updated Objective

## Dataset Creation:

- Generate meaningful datasets for training RL models using AFL.

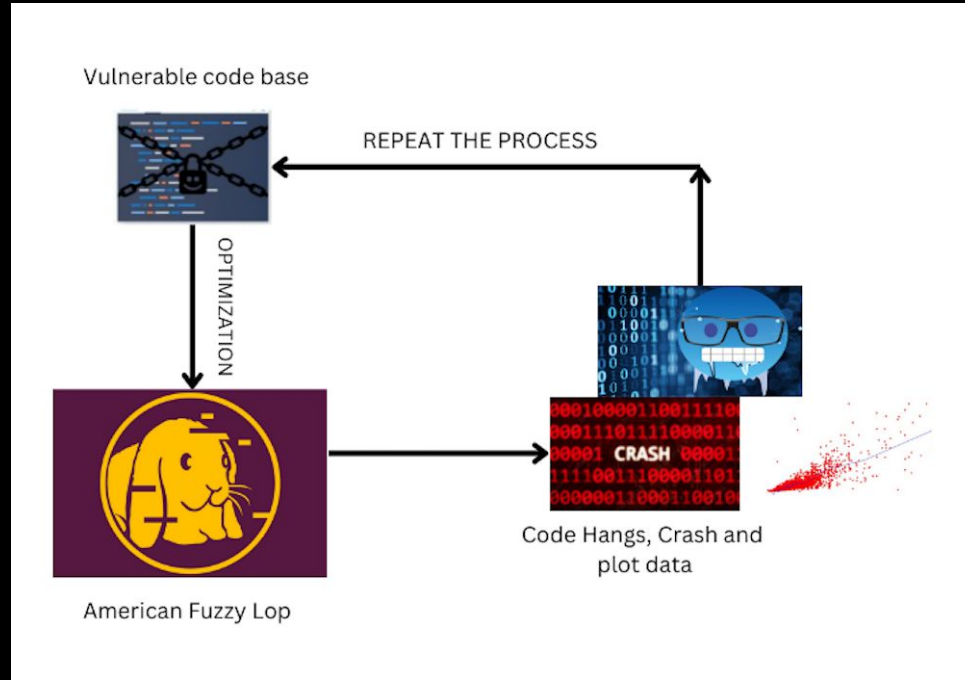
## Optimization:

- Focus on AFL's features to produce proper outputs for RL training.

## Outcome:

- Improve path exploration and crash detection efficiency in fuzz testing.

# Flowchart



# OUTPUT

american fuzzy lop 2.57b (aflrl)			
process timing		overall results	
run time : 0 days, 0 hrs, 0 min, 6 sec		cycles done : 0	
last new path : 0 days, 0 hrs, 0 min, 0 sec		total paths : 109	
last uniq crash : 0 days, 0 hrs, 0 min, 1 sec		uniq crashes : 6	
last uniq hang : none seen yet		uniq hangs : 0	
cycle progress		map coverage	
now processing : 0 (0.00%)		map density : 0.13% / 0.49%	
paths timed out : 0 (0.00%)		count coverage : 1.89 bits/tuple	
stage progress		findings in depth	
now trying : calibration		favored paths : 1 (0.92%)	
stage execs : 0/8 (0.00%)		new edges on : 59 (54.13%)	
total execs : 22.8k		total crashes : 19 (6 unique)	
exec speed : 3115/sec		total tmouts : 0 (0 unique)	
fuzzing strategy yields		path geometry	
bit flips : 13/64, 1/63, 2/61		levels : 2	
byte flips : 0/8, 0/7, 1/5		pending : 109	
arithmetics : 12/444, 0/31, 0/0		pend fav : 1	
known ints : 1/38, 0/196, 0/220		own finds : 107	
dictionary : 0/0, 0/0, 0/0		imported : n/a	
havoc : 0/0, 0/0		stability : 100.00%	
trim : 0.00%/1, 0.00%			
AC		[cpu:287%]	

```
/******  
WARNING: Vulnerability exists below  
  
The line of code below frees the memory block referenced by *top if  
the length of a JSON array is 0. The program attempts to use that memory  
block later in the program.  
  
Diff      - Added: free(*top);  
Payload    - An empty JSON array: []  
Input File - emptyArray  
Triggers   - Use after free in json_value_free()  
*****/  
  
        free(*top);  
/***** END vulnerable code *****/  
  
        break;  
    }  
  
    if (! (value->u.array.values = (json_value **) json_alloc  
        (state, value->u.array.length * sizeof (json_value *), 0)) )  
    {  
        return 0;  
    }
```

```
root@ip-172-31-31-187:/home/ubuntu/aflrl/out# ls  
crashes fuzz_bitmap fuzzer_stats hangs plot_data queue  
root@ip-172-31-31-187:/home/ubuntu/aflrl/out# cd crahes  
bash: cd: crahes: No such file or directory  
root@ip-172-31-31-187:/home/ubuntu/aflrl/out# ls ./crashes  
README.txt id:000002,sig:11,src:000000,op:havoc,rep:8 id:000005,sig:11,src:000000,op:havoc,rep:2  
id:000000,sig:11,src:000000,op:arith8,pos:5,val:-5 id:000003,sig:06,src:000000,op:havoc,rep:8  
id:000001,sig:06,src:000000,op:havoc,rep:8 id:000004,sig:11,src:000000,op:havoc,rep:2
```

## Challenges and Future work

- Environmental and version mismatches
- Transition from RACING to AFLRL required some work.

### Future work

- Train RL with generated crashes.
- Optimize AFL for better vulnerability detection and RL integration.

**Thank You**